



LIBERO-Spatial | ACT vs Diffusion Policy

One suite-conditioned model / seed 0 · same 500 task-init pairs · ACT K=8 chunk; DP DDIM=10, Ta=8

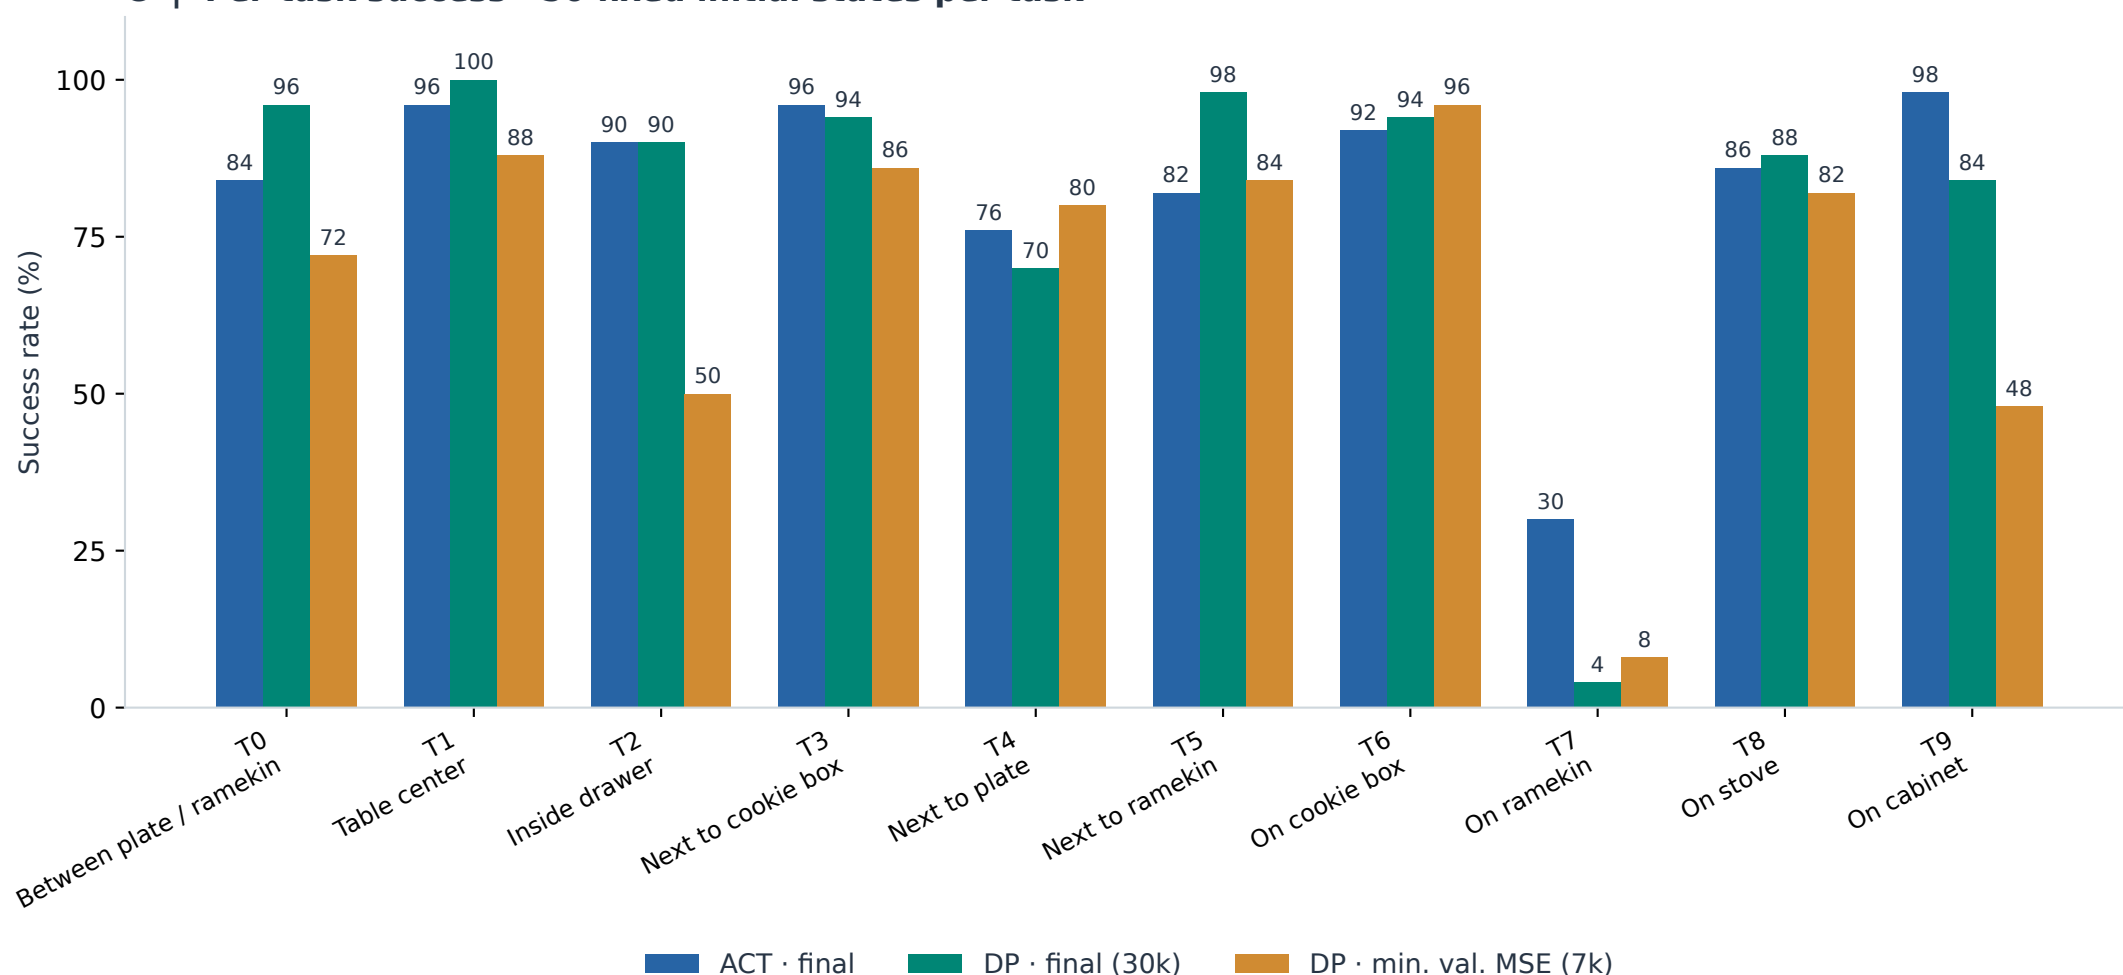
A | Matched evaluation



B | Full chunk inference latency



C | Per-task success · 50 fixed initial states per task



Single training seed; no cross-seed error bars. Training budgets and visual encoders differ.

Latency includes GPU preprocessing, full prediction and action D2H; excludes observation H2D, simulator and network.

DP final was evaluated after observing the 7k result; this is exploratory checkpoint analysis.